SLF Practice Management App

Roadmap Summary — Phase 3 through Phase 6

Phases 0–2 (document intelligence infrastructure, the hearing log and task/reminder engine, and consultation intake with document bundling) are built, tested, and committed. This document summarizes the remaining roadmap — Phase 3 through Phase 6 — as scoped originally, along with the two build decisions already made that affect it.

Decision	Choice
OCR provider (Phase 0, already built)	Local Tesseract only — no cloud OCR vendor
Consultation intake surface (Phase 2, already built)	Staff-assisted only — no client-facing portal yet
Court data source (Phase 5, below)	Licensed third-party legal-data API — not scraping public portals

Phase 3 — Chronology, Synopsis, Remedy Outline

Extends the AI Copilot from single-document summaries to matter-wide analysis

Builds on the existing Copilot's document summarization (already reads case documents via documentTextExtractor.js, now OCR-capable per Phase 0) by extending it across an entire matter instead of one file at a time.

Scope

- A matter-wide **chronology** generator — pulls dated events out of every document/hearing attached to a case and orders them into a timeline.
- An **issue matrix** — the distinct legal issues raised across the matter's documents, cross-referenced to where each appears.
- A **remedy outline** — a structured draft of the relief being sought, built from the chronology and issue matrix rather than typed from scratch.
- Every generated line is **source-linked** back to the specific extracted document and date it came from — consistent with the Copilot's existing design principle of never asserting something it can't point back to.

Depends on

- Phase 0's document text extraction/OCR (the raw material this reads).
- Phase 1's hearing log (structured hearing dates/orders feed the chronology directly, not just document text).

Phase 4 — Drafting Automation

Template-driven drafts, triggered by real dates, never auto-filed

Scope

- A **drafts** table plus a template library, starting with the RFP's named top 5: interlocutory application, adjournment, exemption, delay condonation, and stay.
- **Trigger logic** keyed directly off Phase 1's hearing/task dates — e.g. a hearing logged with a suggested_draft already produces a 'draft_needed' task (built in Phase 1); this phase is what turns that task into an actual generated draft.
- Routed through the **same human-vetting/approval mechanism** the Copilot's propose_change pattern already uses (and that Phase 2's consultation approval mirrors) — a generated draft is never auto-filed or sent; a person reviews and approves it first.

Depends on

- Phase 1's tasks/reminder engine (the trigger source).
- Phase 3's chronology/issue matrix (context a draft is generated from).

Phase 5 — Court Data Sync

Karnataka HC + eCourts — hardest, most fragile; built last on purpose

Flagged as a risk in both of the RFP's source documents: public court portals have no stable API. **Decision made:** this app will integrate through a licensed third-party legal-data API rather than scraping the public HC/eCourts pages directly — trading a vendor cost for reliability and avoiding the ToS/fragility risk of a scraper that breaks every time a court site's layout changes.

Scope

- An **adapter-pattern** integration layer — same shape this app already uses for OCR/e-signature/KMS providers, so the vendor can be swapped later without touching callers.
- **court_sync_logs** — an audit trail of every sync attempt, success or failure.
- **Manual-override controls** — staff can correct a case's synced state by hand when the feed is wrong or stale, without that correction being silently overwritten on the next sync.
- **Cause-list watchlists** — cases flagged so staff are alerted when they appear on an upcoming cause list.

Why last

Built after everything else so the rest of the app already has real, standalone value even if this integration breaks, needs rework, or the vendor relationship changes.

Phase 6 — Notifications, 2FA, Advanced Intelligence

The remaining infrastructure and intelligence work

Scope

- **Notification channels** — email/SMS/WhatsApp Business — nothing is wired today; Phase 1's tasks and hearings are currently surfaced only in-app (dashboard tiles, the task list), by deliberate scope decision made when that phase was built.
- **Two-factor authentication** on login.
- **Limitation watch engine** — tracking statute-of-limitations deadlines with proactive risk flags, extending the risk_flag field Phase 1's hearings table already has.
- **Real Postgres full-text search**, replacing the current simple keyword scorer.
- **Management analytics** — firm-level reporting across cases, billing, and staff workload.
- An **INTAKE_STAFF role**, if the RFP's distinct intake-staff role is wanted separately from the existing PARALEGAL/ASSOCIATE roles.

Depends on

- Phase 1's task/reminder engine (notification channels deliver what it already generates).
- Phase 1's risk_flag column (the limitation watch engine extends it).

Generated as a planning summary for the SLF Document Vault project. Reflects the roadmap as scoped at the end of Phase 2 — subject to change as later phases are actually built.