SLF Practice Management

How the system works, and how it helps lawyers

SLF Practice Management follows a legal matter from the first phone call to the final bill in a single system. It brings together a secure document vault, a court-hearing log, an assistant that reads the firm's own case files, automated first drafts, time-keeping, billing and client-trust accounting, and a client-facing portal — with the access controls, audit trail and approval gates a firm needs before it can safely rely on any of them.

This document explains what the system does, how the parts fit together, what each role in the firm gets out of it, and — just as plainly — what is not connected yet.

1. The problem it addresses

A litigation practice rarely loses time inside any one task. It loses time in the gaps between them: the order sheet that lives in one person's file, the next date that lives in someone's diary, the chronology rebuilt from scratch before every hearing, the application typed again from a precedent someone half remembers, the billable hour reconstructed at month end, and the client who telephones to ask what is happening because there is nowhere for them to look.

> *The system is built around one observation: the firm already holds the information. The work is in finding it, keeping it current, and never missing a date.*

2. The matter lifecycle

The sections below follow a matter in the order a firm actually encounters it.

2.1 First contact — consultation intake

Before any case exists, staff record a prospective client: contact details, the matter category (drawn from a fixed, validated taxonomy of eight practice areas rather than free text), the opposing party, how they were referred, and the intake note itself. A conflict check runs immediately against every party name the firm has ever recorded — including on closed matters — and surfaces as a warning straight away, rather than being discovered later once the workup is already done.

Documents gathered at this stage are uploaded against the consultation and can be merged into a single PDF bundle for the reviewing lawyer, so a partner opens one file rather than eleven. The consultation is then submitted for approval. A partner approves it, at which point it can be converted into a real case, or declines it with a written reason — which returns it to draft so staff know exactly what to fix rather than hitting a dead end.

Intake is deliberately staff-assisted, with a dedicated **Intake Staff** role that can do this work without being handed access to the firm's tasks, hearings or time records.

2.2 Opening the matter — conflict clearance and engagement

Matter Intake walks a new case through a sequence it cannot skip: conflict check, engagement letter drafted, sent, signed, retainer received, activated. A conflict hit places the matter on hold and blocks every later step until a partner clears it as a separate, reasoned decision — never as a side effect of the same request that created the case. The signature step accepts either a countersigned document uploaded by staff or a completed e-signature request. The retainer step posts into real trust accounting, or is explicitly waived for contingency and pro bono work. The case cannot become active until every prior step is genuinely satisfied.

Firms that do not want this ceremony for every matter are not forced into it — the quick “New Case” route still exists and activates immediately.

2.3 The document vault

Every document the firm holds sits in one place, organised by folder, case and tag, and searchable. Uploads are checked three ways: the file extension against an allowlist, the file's actual binary signature against what that extension claims (which is what catches a renamed executable), and a live malware scan. Each file is content-hashed before anything is written, so a byte-identical document already on the matter is caught before a duplicate is created — with the option to keep it anyway when a second copy is genuinely intended.

Re-uploading a file creates a **new version** rather than overwriting the old one. Every version keeps its uploader, timestamp and comment, and two versions of a text-based document can be compared line by line. Scanned orders and photographed papers are made readable by an optical-character-recognition pipeline that runs **on the firm's own machine** — no client document is sent to a cloud OCR vendor. PDFs can be merged and split in place.

2.4 The hearing log and the reminder engine

A hearing is listed; after the appearance an associate records what actually happened — what was submitted, what order was passed, the next date the court gave, a risk flag, and, where relevant, the draft the hearing calls for. That single structured entry then drives the rest of the system on its own:

- a next date produces a **prepare-for-next-hearing** task, dated two days ahead of the listing, so preparation is not discovered on the morning itself;
- a named draft produces a **draft-needed** task, due on the next date, which is what later becomes an actual generated draft;
- a hearing whose date has passed with no update recorded appears on the dashboard as **not updated after hearing** — the practice-management equivalent of an unreconciled entry.

These follow-ups are re-derived every time the update is saved, so correcting a typing error refreshes the existing task rather than multiplying the task list.

2.5 Matter intelligence

An assistant sits over the firm's own files and answers questions about them in ordinary language — finding documents, summarising them from their actual contents rather than their filenames, and pulling together case details across a matter. Beyond single documents, the system produces three matter-wide work products:

- **Chronology** — dated events extracted from every document and logged hearing on the matter, ordered into a single timeline, with conflicting dates flagged rather than silently resolved;
- **Issue matrix** — the distinct legal issues raised across the matter, each cross-referenced to every place it appears;
- **Remedy outline** — a structured draft of the relief sought, built from the chronology and issue matrix rather than typed from scratch.

Every generated line is source-linked back to the specific document or hearing it came from. Any citation the model produces that does not correspond to a real document or hearing on that matter is discarded before it ever reaches the screen — a fabricated reference cannot be dressed up as a real one.

2.6 Drafting automation

The draft-needed task becomes an actual draft. Five templates cover the applications a litigation practice files most often — interlocutory application, adjournment, exemption, delay condonation, and stay. Each is generated from the matter's own chronology and issue matrix (producing them first if they do not exist yet), follows that template's required sections, and carries the same source links back to the underlying material.

Nothing is ever auto-filed or sent. A generated draft is created in a pending-review state, and only a partner can approve or reject it. Approval closes the originating task; rejection requires a written reason and reopens the task so the draft can be regenerated. Each generation is kept rather than overwritten, so a reviewer can see what changed between attempts.

2.7 Court data sync

Case status is synchronised through a licensed legal-data provider rather than scraped from the public High Court and eCourts pages — a deliberate decision that accepts a vendor cost in exchange for reliability, and avoids both the terms-of-service exposure and the constant breakage of a scraper that fails whenever a court website is redesigned. The integration sits behind an adapter, so the provider can be changed later without touching anything that depends on it.

Every sync attempt, successful or failed, is recorded in an audit trail. Where a feed is wrong or stale, staff can correct the case's position by hand; doing so places the matter into **manual override**, and the next sync refuses to run rather than silently overwriting the correction — it has to be released deliberately. Matters can also be flagged to a cause-list watchlist so staff are alerted when one is due on an upcoming list.

2.8 Limitation watch

Statute-of-limitations deadlines are tracked against each matter with a risk flag that escalates on its own as the date approaches: ordinary beyond forty-five days, **watch** within forty-five, and **urgent** within fourteen days or already lapsed. The flag is recalculated every time it is read rather than stored, so it can never quietly go stale because nobody happened to open that record. Resolving a deadline records who closed it, when, and why.

2.9 Time, billing and client trust money

Time is recorded against a matter in whole minutes, with the fee-earner's rate captured on the entry at the moment it is logged — so a later change of rate never retroactively alters what an already-billed hour claims to have been worth. Unbilled entries from a single case roll up into a properly itemised invoice in one action, and the invoice, its line items and each entry's status all commit together or not at all. Once an entry has been invoiced it becomes immutable: a correction is made by a new entry, never by rewriting what was already billed.

Client trust money is kept structurally separate from firm revenue rather than as another column on the same table. The ledger is append-only — a correction is always a new reversing entry, never an edit or a deletion. A deposit or disbursement that would take a client's balance negative is refused, and one client's money can never be spent on another client's matter within the same pooled account. The only route from trust into firm revenue is a transfer to operating, which marks the linked invoice paid in the very same database transaction as the withdrawal, so a retainer can never appear drawn down while the invoice it paid still shows unpaid. A three-way reconciliation — bank statement against ledger balance against the sum of individual client balances — can be run at any time and saved as a dated snapshot.

> *Every trust route, reads included, is restricted to partners, and the assistant has no access to any of it whatsoever.*

2.10 The client portal and communications

Clients sign in to their own simplified portal showing their matters, documents, messages and invoices — and nothing else. Messages are threaded between staff and client, and internal staff notes on the same thread are removed from the response before it leaves the server, not merely hidden in the interface. Attachments are real references to vault documents, re-checked against the sender's own permissions each time one is posted. Signature requests are issued as one-time links to the firm's own signing page.

3. What each role gets

Role	What the system does for them
Partner / Senior Partner	The firm-level picture: revenue to date, overdue invoices, matters by status and practice area, and each fee-earner's open tasks and hours this month. Holds the decisions that should require judgement — clearing conflicts, approving consultations, approving or rejecting generated drafts, marking invoices paid, and every movement of client trust money.
Associate	Records the hearing update once and receives the follow-up tasks, the draft and the refreshed chronology from it. Generates matter-wide analysis and first drafts instead of assembling them by hand, and works from a timeline that is already sourced back to the underlying documents.
Paralegal	Runs the day-to-day: listing hearings, working the task list, tracking limitation dates, linking and syncing court records, sending notifications and logging their own time — without being given access to the firm's financial controls.
Intake staff	A dedicated role for booking and working up prospective clients: intake notes, conflict warnings, document bundles and submission for approval, with no access to the firm's tasks, hearings or time records.
Administrator	Provisions users by invitation only — there is no self-signup anywhere in the system — ends every session of a departing employee at once, reads the complete audit log, and manages the firm's storage quota.
Client	Sees their own matters, documents, messages and invoices, and nothing belonging to anyone else. That restriction is applied in the database query itself, not by hiding things in the interface.

4. Why a practitioner can rely on it

Software that touches privileged material, court deadlines and client money has to earn its place. Six principles run through the system:

Nothing generated is asserted without a source

Every chronology entry, identified issue and drafted section points back to the specific document or hearing behind it. Citations that do not match real material on that matter are filtered out before display, so an invented reference cannot pass as a real one.

Nothing generated acts on its own

The assistant's tools are read-only. A request to change something comes back as a proposal for a person to approve, and the conversational loop itself cannot write data even if the model attempts it. Billing and trust accounting have no assistant tool at all — that exclusion is built into the structure of the system, not merely written into an instruction the model is asked to follow.

Nothing reaches a court without a person

Generated drafts require partner approval and are never filed or sent automatically. The same human-approval pattern governs consultations, conflict clearance and case changes proposed by the assistant.

Access is enforced where it counts

A client's restriction to their own matters is applied server-side on every request. The automated test suite actively attempts the forbidden reads — one client trying to reach another's case, document, invoice or message — and requires them to fail.

Everything consequential is written down

Each creation, amendment, approval, sync attempt, notification and trust movement writes an audit entry naming who did it and when. The audit log is readable by administrators and is never edited by the application.

Corrections add; they do not erase

Trust entries, invoiced time and document versions are all append-only. A mistake is corrected by recording a further entry, so the record of what actually happened survives — which is precisely what a firm needs when a file is examined months later.

Credentials and content are protected

Two-factor authentication works with any standard authenticator application and issues single-use recovery codes for a lost device. Sessions can be ended individually or all at once, immediately and everywhere. Client message contents are encrypted at rest under a managed key, so access to the database alone does not yield a readable message. Passwords are scored for genuine strength rather than merely checked for length.

5. The technical foundation, briefly

The system runs on Node.js and PostgreSQL, with thirty-one versioned database migrations and a suite of just over four hundred automated tests across thirty-nine files — including concurrency tests that drive real database races to prove, for instance, that two simultaneous attempts cannot double-bill the same hour or draw the same retainer twice.

Every request carries a correlation identifier through the logs and back to the browser, so a problem reported by a user can be matched to its exact server records. Authentication endpoints are rate-limited. Deployment is containerised with a health endpoint, and the test suite runs on every change.

Each external dependency — optical character recognition, e-signature, key management, court data and notification delivery — sits behind its own adapter. Changing any one vendor means adding a single file, not reworking the code that depends on it.

6. What is not connected yet

Stated plainly, because a firm should know where the edges are before it relies on something:

- **No notification vendor is connected.** Email, SMS and WhatsApp channels are built, audited and ready, but until a provider is configured every send is recorded rather than actually delivered. The system reports this honestly instead of showing a delivery that never happened.
- **No court-data feed is contracted.** The sync layer, audit trail, manual override and cause-list watchlist are all built and working; a licensed provider has to be connected before a sync returns live data. Manual override covers the interim.
- **No third-party e-signature provider.** Signing works end to end on the firm's own page today; connecting DocuSign or a similar service is one additional adapter file.
- **No email delivery.** An invitation link is shown once on screen for the administrator to pass to the new joiner by whatever means they already use.
- **Document search covers metadata, not full contents.** Filenames, matter names and tags are fully searchable; searching the text inside every document would require storing extracted text at upload.
- **The client portal has no self-service.** Clients cannot register, reset their own password or receive alerts; they see the current position whenever they open a page.

Each of these is a documented boundary and, in most cases, a commercial decision for the firm to make — not engineering work abandoned half-finished.

Prepared as a working overview of the SLF Practice Management system. Descriptions reflect the system as built and tested; the boundaries in section 6 are accurate as at the date of this document.